

Cosine Similarity Benchmark

Eric Crockett

May 9, 2025

List of Corrections

Eric: Feel free to improve the variable names	1
Eric: I suspect Shai had a different algorithm in mind since he said it may take up to depth 8 for this step.	1
Eric: Is it possible to use scheme-switching to do a bit-by-bit comparison?	3

1 Introduction

This document describes Shai Halevi’s algorithm for cosine similarity. We have a dataset with N (key, value) records. Each key is an ℓ_2 -normalized m -dimensional vector. Let A be the $m \times N$ matrix with all of these keys. We assume that each value fits in k slots.

There is a single secret query (row) vector v of dimension m , also ℓ_2 -normalized.

There is a threshold $\tau < 1$, and a promise that no more than M records in the dataset are τ -close to v .

The goal is to fetch the values corresponding to the keys that are τ -close to v .

The server has encryptions of the key and value for each record. Shai assumes we are using a ring dimension of size 65K, so there are 32K slots (note that K is never a variable). The encrypted data is batched into blocks of size 32K records, where the keys for block i are encoded into the slots of m ciphertexts, and the values for block i are encoded into the slots of k additional ciphertexts. The number of blocks is $n = \lceil N/32K \rceil$.

The following list summarizes the notation:

- N the number of records
- m the number of slots per record key
- k the number of slots per record value
- $n := \lceil N/32K \rceil$, the number of ciphertexts required to hold N slots
- τ the dot product threshold. We return values for keys whose dot product with the query is $\geq \tau$.
- M the maximum number of keys above the dot product.
- C chunk size, used in the compaction step.
- $r := \lceil \log_{\frac{M \cdot C}{32K}}(2^{-32}) \rceil$, the maximum number of 1s in a slice of size $C \cdot n$, with probability 2^{-32} .
- $R := \lfloor 32K/C \rfloor$, the number of output ciphertexts

We require $C \cdot n \leq 32K$, and $r \cdot k \leq 32K$. Eric: Feel free to improve the variable names

Eric!

The procedure proceeds as follows:

1. **Dot Products** Compute $u = v \cdot A$. We can isolate and duplicate each coefficient of v using a multiplication by a (public) mask, followed by 16 shifts and 15 additions. Then we multiply each coefficient ciphertext by a packed row of A , and add the products together (for each block). This process has depth 2. Eric: I suspect Shai had a different algorithm in mind since he said it may take up to depth 8 for this step. See DotProds in Algorithm 1. Eric!
2. **Thresholding** Compare the slots of u against the threshold τ . Shai suggests using a polynomial approximation for this step, consuming 6 or 7 levels.
3. **Compacting** The algorithm requires a promise that at most M dot products will be larger than τ , so the comparison vector has at most M 1s in it. The output is R ciphertexts, each with at most $r \cdot k$ non-zero left-justified slots. Roughly, the algorithm involves “slicing” the threshold vector and value matrix by taking C slots from each of the n ciphertexts (in the input, or in one row of the value matrix). We compact these slots into a single ciphertext, then compute a running sum. With high probability, the sum will be $\leq r$, by design. For $i \in [1, r]$, we do an equality test on the running sum and multiply by the threshold vector to obtain an indicator vector with a one in the position of the i^{th} one in the threshold vector. We can then extract the corresponding value coefficient from the value matrix, and compactly store it in an output ciphertext. See Algorithm 2 for details.

Algorithm 1 Encrypted vector/matrix multiplication. v is a single CT with data in first m non-zero slots. A is a matrix with m rows, where each row is composed of n ciphertexts. The j^{th} ciphertext in the i^{th} row is $A_{i,j}$.

```

1: procedure REPLICATE( $x$ )                                ▷  $x$  has a value in one slot and 0s elsewhere
2:                                                         ▷ Output has this value in all 32K slots
3:   for  $i$  in range(15) do                                ▷ Hard-coded for 32K slots
4:      $x \leftarrow x + (x \gg 2^i)$ 
5:   return  $x$ 
6:
7: procedure EXTRACTCOEFF( $v, i$ )                          ▷ Return a CT with  $v_i$  in all slots
8:    $mask \leftarrow \text{Indicator}(i)$ 
9:    $w \leftarrow mask \cdot v_0$ 
10:  return Replicate( $w$ )
11:
12: procedure DOTPRODS( $v, A$ )
13:  prods  $\leftarrow [0] * n$ 
14:  for  $j$  in range( $m$ ) do
15:     $v_i \leftarrow \text{ExtractCoeff}(v, j)$ 
16:    for  $i$  in range( $n$ ) do
17:      prods[ $i$ ]  $\leftarrow$  prods[ $i$ ] +  $v_i \cdot A_{i,j}$ 
18:  return prods

```

Algorithm 2 Compaction. us is the output of the thresholding step, a vector with n ciphertexts. ps is a $k \times n$ matrix where each row ps_i is n packed ciphertexts. $ps_{i,j}$ means the j^{th} ciphertext in the i^{th} row

```

1: procedure RUNNINGSUM( $v$ )
2:    $rs \leftarrow v \cdot \text{Indicator}(0)$ 
3:    $acc \leftarrow v$ 
4:   for  $i$  in range(1, 32K) do                                ▷ Extracts the  $i^{\text{th}}$  component of the running sum
5:      $acc \leftarrow v + (acc \gg 1)$ 
6:      $rs \leftarrow rs + (acc \cdot \text{Indicator}(i))$ 
7:   return  $rs$ 
8:
9: procedure EXTRACTSLICE( $xs, t$ )                            ▷  $xs$  is a list of  $n$  CTs,  $t \in [0, R)$ 
10:   $mask \leftarrow [1] * C + [0] * (32768 - C)$ 
11:   $x \leftarrow 0$ 
12:  for  $i$  in range( $n$ ) do
13:     $x \leftarrow x + (((xs_i \cdot (mask \gg C \cdot t)) \ll C \cdot t) \gg C \cdot i)$ 
14:  return  $x$ 
15:
16: procedure COMPACTIFY( $us, ps$ )
17:  accs  $\leftarrow [0] * R$ 
18:  for  $t$  in range( $R$ ) do                                ▷ Loop over each slice
19:     $u \leftarrow \text{ExtractSlice}(us, t)$ 
20:     $v \leftarrow \text{RunningSum}(u)$ 
21:     $xs \leftarrow [0] * r$ 
22:    for  $i$  in range(1,  $r + 1$ ) do
23:       $w_i \leftarrow \text{map}((= i), v)$                                 ▷ 1 where the running-sum is  $i$ , 0 elsewhere
24:       $xs_i \leftarrow w_i \cdot u$                                 ▷ Extracts the  $i^{\text{th}}$  1 from  $u$ 
25:      for  $i$  in range( $k$ ) do                                ▷  $k$  is the number of slots per value
26:         $p \leftarrow \text{ExtractSlice}(ps_i, t)$ 
27:        for  $j$  in range(1,  $r + 1$ ) do                                ▷ We assume that there are at most  $r$  1s in  $u$ 
28:           $x \leftarrow p \cdot xs_j$                                 ▷ The value corresponding to the  $j^{\text{th}}$  1 from  $u$ 
29:           $y \leftarrow \text{Replicate}(x)$                                 ▷ The value is now in all slots
30:           $z \leftarrow y \cdot \text{Indicator}(j \cdot k + i - 1)$         ▷ The value is now in the appropriate output slot
31:          accs[ $t$ ]  $\leftarrow$  accs[ $t$ ] +  $z$ 
32:  return accs

```

2 Failure Modes

2.1 Broken Promise

If there are in fact $> M$ matches, this does not necessarily invalidate the results. The effect is more indirect in that it increases the probability that more than r keys match in a given slice.

2.2 Unlucky Distribution

The algorithm returns up to r results per slice, and therefore up to $r \cdot R$ results total. If more than r keys match for a given slice (either because we got *very* unlucky, or because there are more than M matches and we are only slightly unlucky), only up to r results are returned per slice.

2.3 High-Precision Threshold

One problem that Shai identified with the thresholding step is that the polynomial approximation will return values ≈ 1 if the dot product is larger than τ and ≈ 0 if the dot product is smaller than τ . However, for values very close to τ , the polynomial approximation will return values near $1/2$, which results in inaccurate results. One way to prevent this is to require a “promise” that the dot products are sufficiently far away from τ . **Eric: Is it possible to use scheme-switching to do a bit-by-bit comparison?**

Eric!